



FlashMLA Decode — Batch Scaling (L = 32K)

Ours vs Baseline


Kernel Duration



Aggregate Throughput



Effective DRAM BW



Speedup vs Baseline

